

Hint for OS HW9

Write a kernel module and a user-mode program to show how the logical address of a variable in the user-mode program is converted to the physical address in x86-64 environment.

All Info from here -> [Intel Software Developer Manual](#)

Step 0 How to build a Kernel module ?

FYI. [EN](#) / [CN](#)

Step 1 Virtual Memory Address to Linear Address

In Manual P83:

3.7.3.1 Memory Operands in 64-Bit Mode

In 64-bit mode, a memory operand can be referenced by a segment selector and an offset. The offset can be 16 bits, 32 bits or 64 bits (see Figure 3-10).




Figure 3-10. Memory Operand Address in 64-Bit Mode

And since the value is saved on the stack, it seems we need to get the segment selector:

Manual P161

Data manipulation instructions that specify the EBP register as a base register automatically address locations within the stack segment instead of the data segment.

```
0000000000000068a <main>:
68a: 55                push  %rbp
68b: 48 89 e5          mov   %rsp,%rbp
68e: 48 83 ec 10       sub   $0x10,%rsp
692: 64 48 8b 04 25 28 00 mov   %fs:0x28,%rax
699: 00 00
69b: 48 89 45 f8       mov   %rax,-0x8(%rbp)
69f: 31 c0            xor   %eax,%eax
6a1: 48 b8 dd cc bb aa 78 movabs $0x12345678aabbccdd,%rax
6a8: 56 34 12
6ab: 48 89 45 f0       mov   %rax,-0x10(%rbp)
6af: b8 00 00 00 00    mov   $0x0,%eax
6b4: e8 97 fe ff ff    callq 550 <getpid@plt>
6b9: 89 c1            mov   %eax,%ecx
6bb: 48 8d 45 f0       lea   -0x10(%rbp),%rax
6bf: 48 89 c2          mov   %rax,%rdx
6c2: 89 ce            mov   %ecx,%esi
6c4: 48 8d 3d 9d 00 00 00 lea   0x9d(%rip),%rdi    # 768 <_IO_stdin_used+0x8>
6cb: b8 00 00 00 00    mov   $0x0,%eax
6d0: e8 8b fe ff ff    callq 560 <printf@plt>
6d5: eb fe            jmp   6d5 <main+0x4b>
6d7: 66 0f 1f 84 00 00 00 nopw  0x0(%rax,%rax,1)
6de: 00 00
```

But in Linux, all processes running in User Mode use the same pair of segments to address instructions and data. In Linux 4.15:

```
// In arch/x86/kernel/process_64.h
```

```

void
start_thread(struct pt_regs *regs, unsigned long new_ip, unsigned long new_sp)
{
    start_thread_common(regs, new_ip, new_sp,
                        __USER_CS, __USER_DS, 0);
}
static void
start_thread_common(struct pt_regs *regs, unsigned long new_ip,
                    unsigned long new_sp,
                    unsigned int _cs, unsigned int _ss, unsigned int _ds)
{
    ...

    regs->ip      = new_ip;
    regs->sp      = new_sp;
    regs->cs      = _cs;
    regs->ss      = _ss;
    regs->flags    = X86_EFLAGS_IF;
    force_iret();
}

// In arch/x86/include/asm/segment.h
#define __KERNEL_CS      (GDT_ENTRY_KERNEL_CS*8)
#define __KERNEL_DS      (GDT_ENTRY_KERNEL_DS*8)
#define __USER_DS        (GDT_ENTRY_DEFAULT_USER_DS*8 + 3)
#define __USER_CS        (GDT_ENTRY_DEFAULT_USER_CS*8 + 3)

```

So, we don't need to pass value of ss in user mode, we can simply use macro `__USER_DS` to get the segment selector. The format of a segment register looks like following: (From Chapter 3, "Protected-Mode Memory Management," in the Intel® 64 and IA-32 Architectures Software Developer's Manual, Volume 3A.)

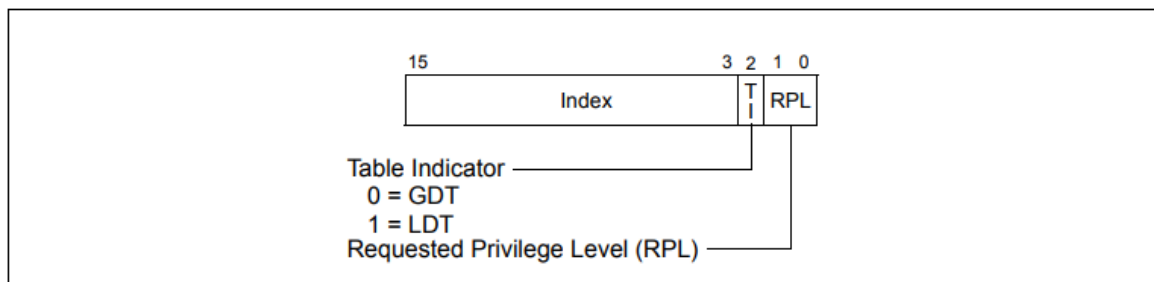
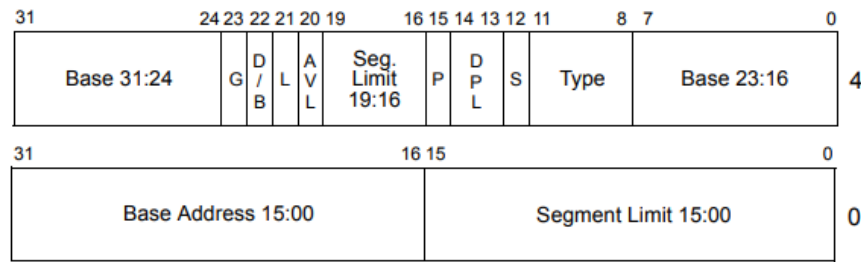


Figure 3-6. Segment Selector

OK, now you get the value of the segment register and its format, the next step you need to do, is get the address of the Global/Local Descriptor Table, to get the descriptor and linear address.

The format of the descriptor.



- L — 64-bit code segment (IA-32e mode only)
- AVL — Available for use by system software
- BASE — Segment base address
- D/B — Default operation size (0 = 16-bit segment; 1 = 32-bit segment)
- DPL — Descriptor privilege level
- G — Granularity
- LIMIT — Segment Limit
- P — Segment present
- S — Descriptor type (0 = system; 1 = code or data)
- TYPE — Segment type

Figure 3-8. Segment Descriptor

Hit: You can use inline assemble to get [G | L]DT, and use desc_ptr struct can make things easier.

```
struct desc_ptr xdr;  
asm("s[g|l]dt %w0":"=m"(xdr));
```

Step 2 Linear Address to Physical Address

Since we have a logical address, if we want to convert it to physical address, we need the information of its pgd. But since we have in Kernel, the value in CR3 is not same as our user program, luckily, we can get it from task struct using pid. But, using this method we get a virtual memory.

Hit: using `virt_to_phys()` to convert.

```
// In include/linux/sched.h
struct task_struct {
    ...
    struct mm_struct      *mm;
    ...
}

// In include/linux/mm_types.h
struct mm_struct {
    ...
    pgd_t * pgd;
    ...
}
```

Now, you get the physical address of pgd, but you can not directly access this address.

Hit: Using `kmap()` to map a frame, don't forget to use `kunmap` to free the space.

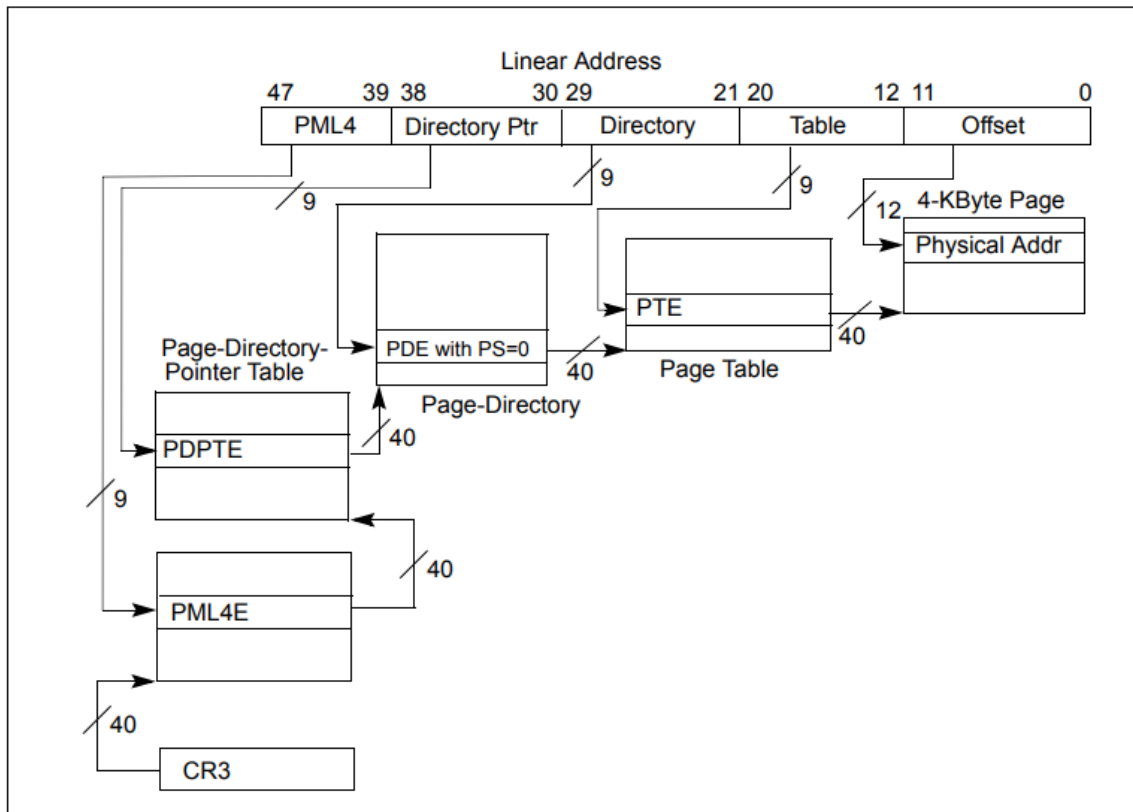


Figure 4-8. Linear-Address Translation to a 4-KByte Page using IA-32e Paging

Next, use the offsets to find the value we want.

Any Problems @ me in QQ groups :)