



File System Implementation

Yajin Zhou (<http://yajin.org>)

Zhejiang University




Two different aspects of FS

- (on-disk) Structure about FS
- Access methods
 - how does the FS maps the calls (open/read/write) onto its structures

An Example

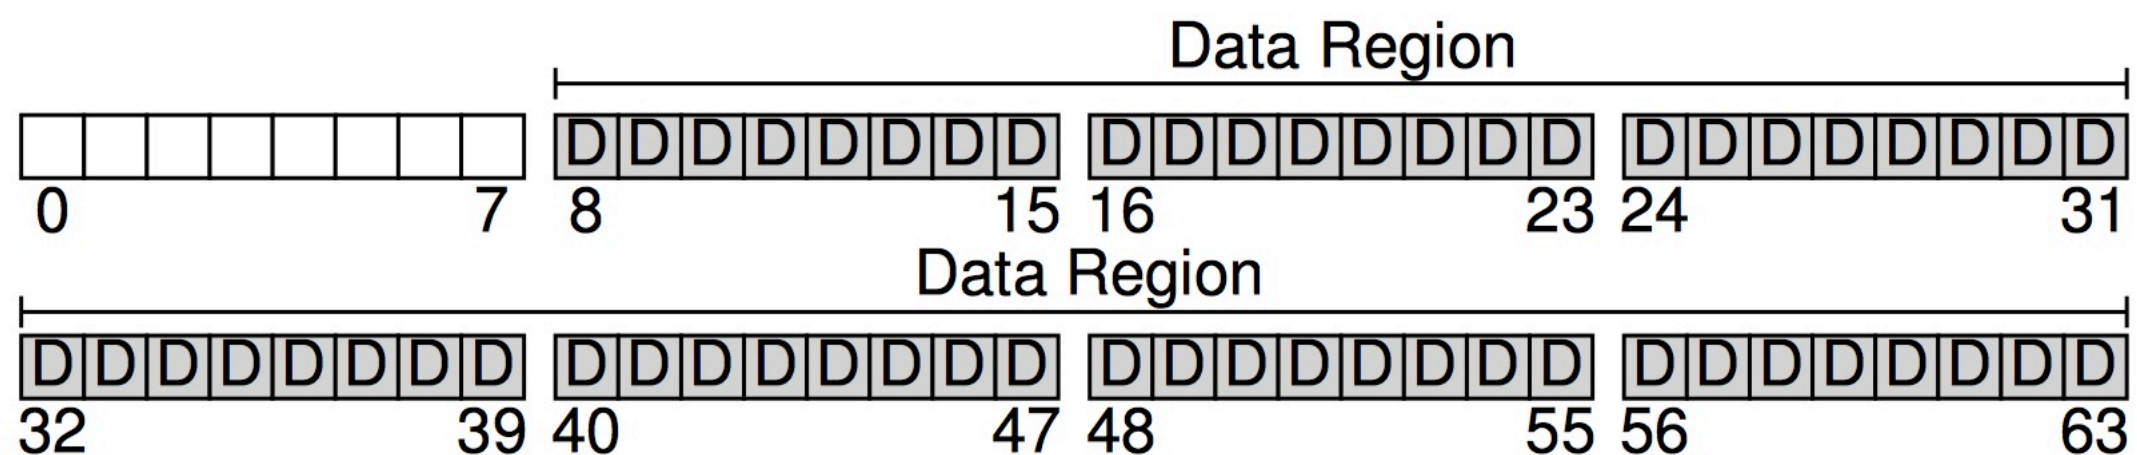
- Suppose we have a serial of blocks
 - Block size: 4k
 - 64 blocks





Data Region

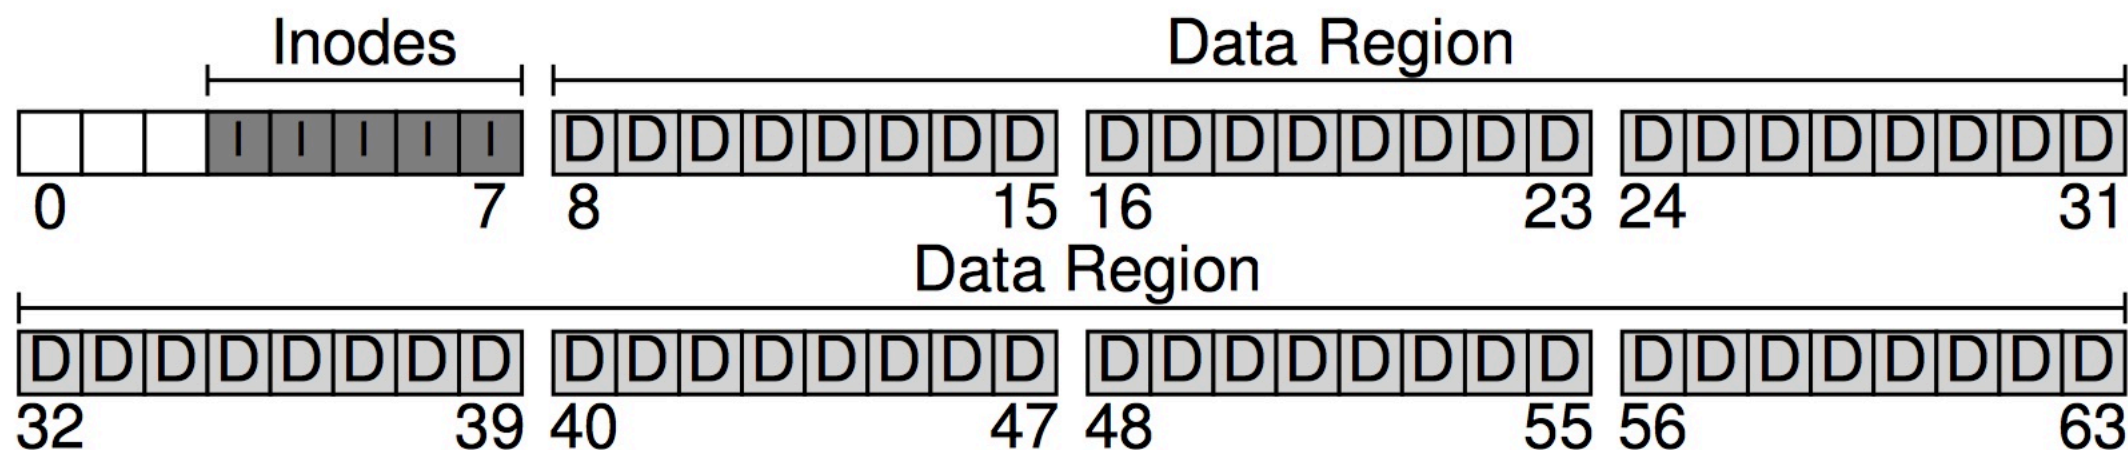
- We reserve some blocks for data
 - 56 of 64 blocks





Inodes

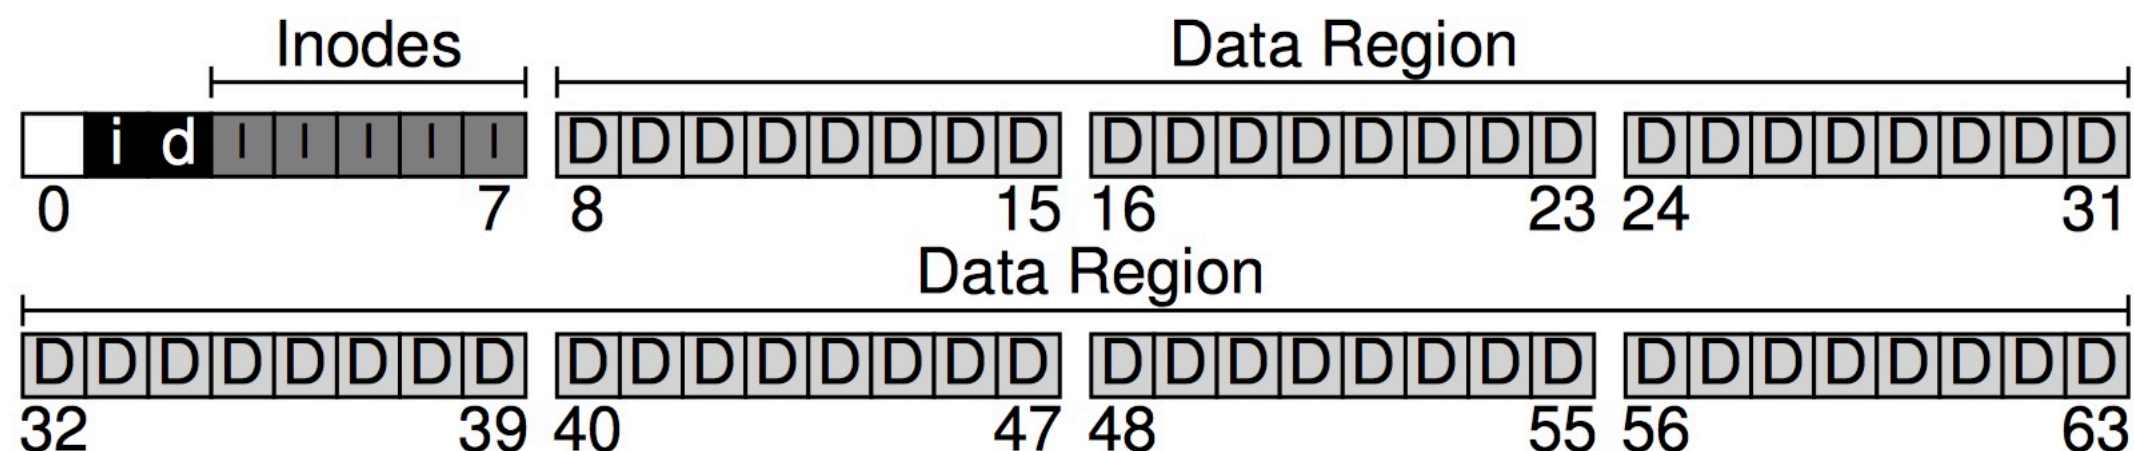
- Inode table: contains inodes
- 5 of 64 blocks are reserved for inodes
- Suppose inodes are 256 bytes, 4 kb block can hold 16 inodes, then 5 blocks -> 80 inodes -> 80 files (directories)



D: Data block
I: inode

Bitmap

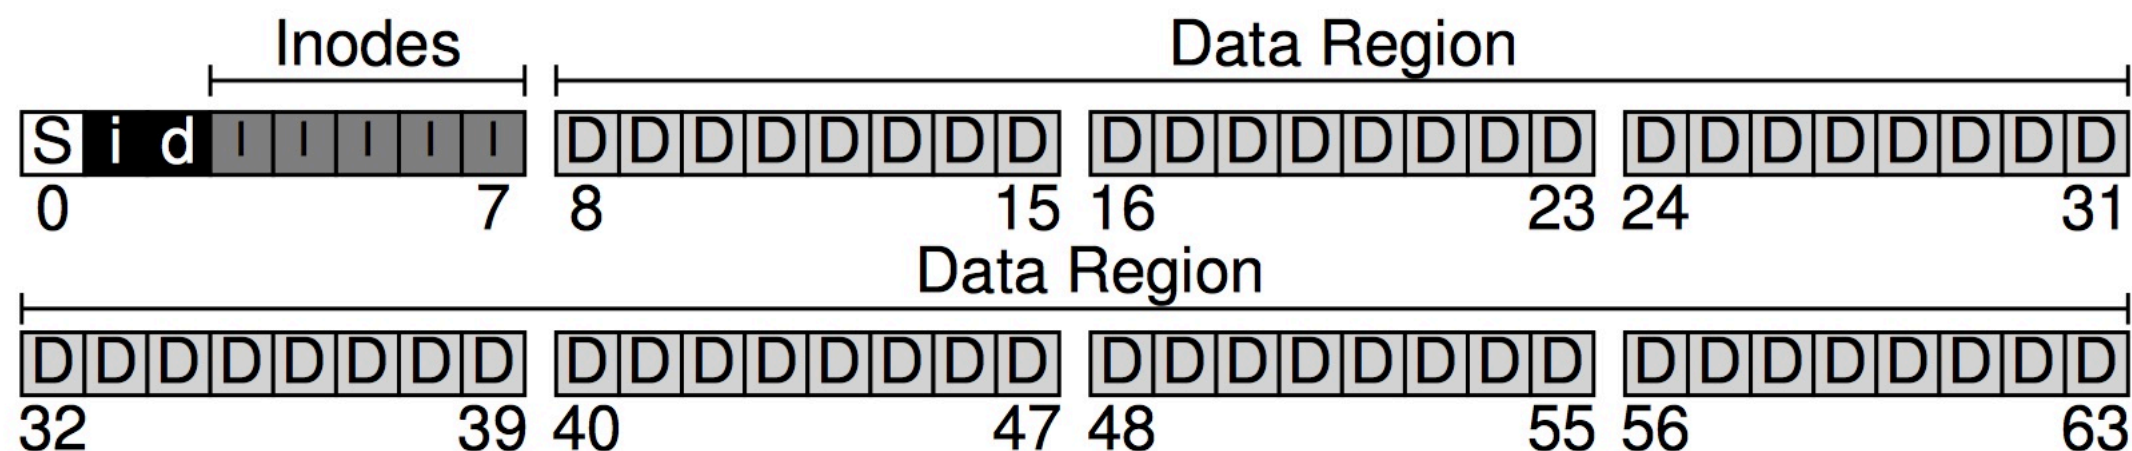
- Suppose we use bitmap to manage the free space
 - One bitmap for free inodes
 - One bitmap for free data region



D: Data block
 I: inode
 I:inode bitmap
 d: data region bitmap

Superblock

- Superblock
- Contains information about this file system: how many inodes/ data blocks, where the inode table begins, where the data region begins, and **a magic number**



D: Data block
 I: inode
 I:inode bitmap
 d: data region bitmap
 S: superblock



Inode

- Each inode is identified by a number(inode number)
- To read inode number 32
 - $32 * \text{sizeof(inode)} = 8k$
 - Address: $8k + 4k(\text{super block}) + 8Kk(\text{BITMAP}) = 20K$

The Inode Table (Closeup)

				iblock 0				iblock 1				iblock 2				iblock 3				iblock 4			
Super	i-bmap d-bmap			0	1	2	3	16	17	18	19	32	33	34	35	48	49	50	51	64	65	66	67
				4	5	6	7	20	21	22	23	36	37	38	39	52	53	54	55	68	69	70	71
				8	9	10	11	24	25	26	27	40	41	42	43	56	57	58	59	72	73	74	75
				12	13	14	15	28	29	30	31	44	45	46	47	60	61	62	63	76	77	78	79
0KB	4KB	8KB	12KB	16KB				20KB				24KB				28KB				32KB			

Ext2 Inode

Size	Name	What is this inode field for?
2	mode	can this file be read/written/executed?
2	uid	who owns this file?
4	size	how many bytes are in this file?
4	time	what time was this file last accessed?
4	ctime	what time was this file created?
4	mtime	what time was this file last modified?
4	dtime	what time was this inode deleted?
2	gid	which group does this file belong to?
2	links_count	how many hard links are there to this file?
4	blocks	how many blocks have been allocated to this file?
4	flags	how should ext2 use this inode?
4	osd1	an OS-dependent field
60	block	a set of disk pointers (15 total)
4	generation	file version (used by NFS)
4	file_acl	a new permissions model beyond mode bits
4	dir_acl	called access control lists

Figure 40.1: **Simplified Ext2 Inode**



Multi-Level Index

- To support bigger file, we need multi-level index for the nodes



Directory Organization

- Suppose a dir (inode number 5) has 3 files: foo,bar, footer
- Name:
- Strlen: length of the name
- Reclen: length of the name plus **left over space (what's this?)**
 - For reuse the entry purpose

inum		reclen		strlen		name
5		4		2		.
2		4		3		..
12		4		4		foo
13		4		4		bar
24		8		7		foobar



Free Space Management

- Bit map
- Some OS will use the pre-allocation policy
 - For instance, when a file is created, a sequence of blocks (say 8) will be allocated
 - This can guarantee that the file on the disk is contiguous



Read /foo/bar

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
open(bar)			read	read	read	read	read			
read()					read		read			
read()					write					
read()					read				read	
read()					write					
					read					
					write					

What about the system-wide/per-process open file table?



Write to Disk: /foo/bar

	data bitmap	inode bitmap	root inode	foo inode	bar inode	root data	foo data	bar data[0]	bar data[1]	bar data[2]
create (/foo/bar)		read write	read	read	read write	read	read write			
write()	read write				read write			write		
write()	read write				write read				write	
write()	read write				write read					write



Caching and Buffering

- Without caching, each file open would require two reads for each level of the directory
 - One for the inode, and one for data
- Early system allocate a **fixed-size** cache to hold popular blocks
- Modern systems use a **unified page cache** for both virtual memory pages and file system pages
- Write buffering: does not write to disk immediately, instead sync to disk for like 5 - 30 seconds
- Database: direct IO with raw data

Log-structured File Systems



Motivation

- System memories are growing:
 - can cache more data, disk operations are mostly write since read are serviced by the cache -> need to optimize write performance
- There is a large gap between random I/O and sequential I/O performance
 - Use the disk in sequential manner

Idea: try to make use of the sequential bandwidth of the disk




LFS

- LFS: log-structured File System
- When writing to disk, LFS first buffers all updates (including metadata) into a **memory segment**; when the segment is full, it is written to disk in **one long and sequential transfer** to an **unused part** of the disk
- LFS **never overwrites existing data**, but rather always writes segments to free locations

Writing To Disk Sequentially

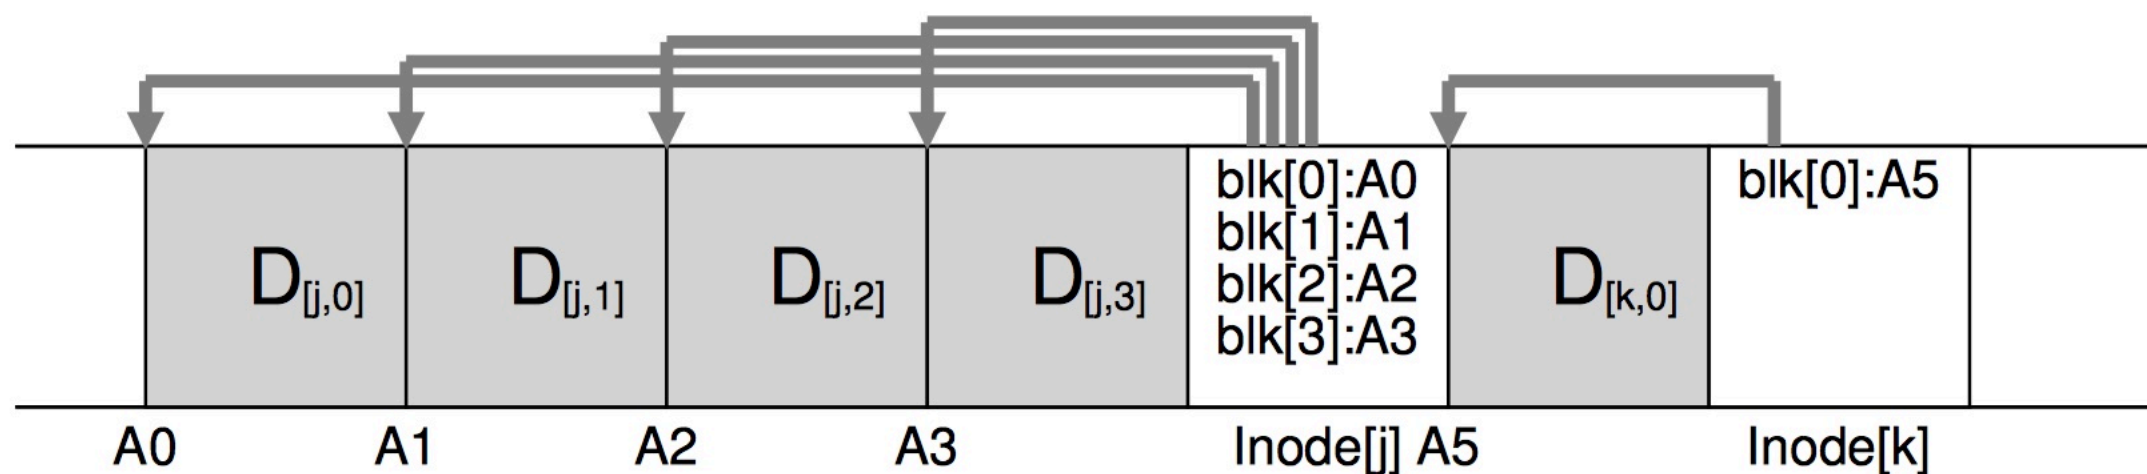
- Write the data block and metadata into the disk
 - I: Inode





Write Buffering

- We can write to the disk when the **memory segment** is full
 - First is writing four blocks to file j
 - Second is one block being added to file k



How to find inode?



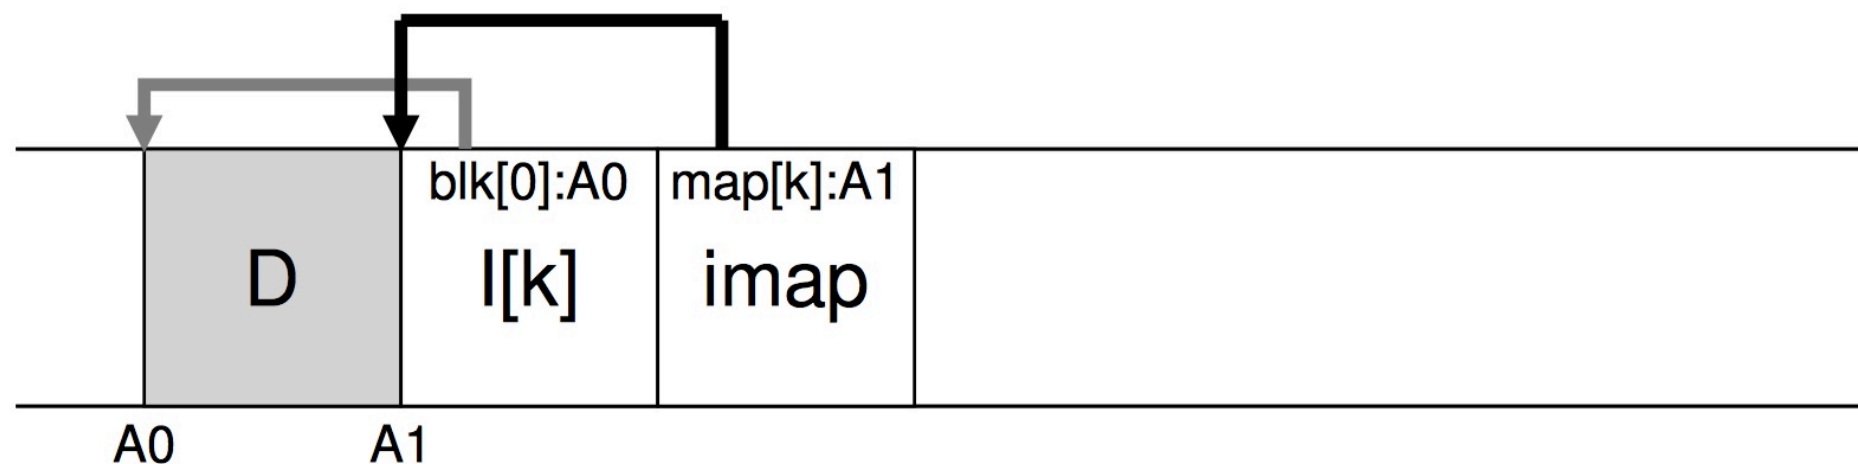
Find Inode

- UNIX FS
 - Keep Inode in fixed locations
 - LFS: is hard
 - Inodes are scattered throughout the disk



Inode Map


- Inode map(imap)
- This map takes an inode number as input and produces the disk address of the most recent version of the inode
- LFS places inode map right next to where it is writing all of the other new information



How to find imap?

Checkpoint Region

- Checkpoint region (CR)
 - Contains pointers to the latest pieces of the inode map
 - CR is updated periodically (say 30 seconds)



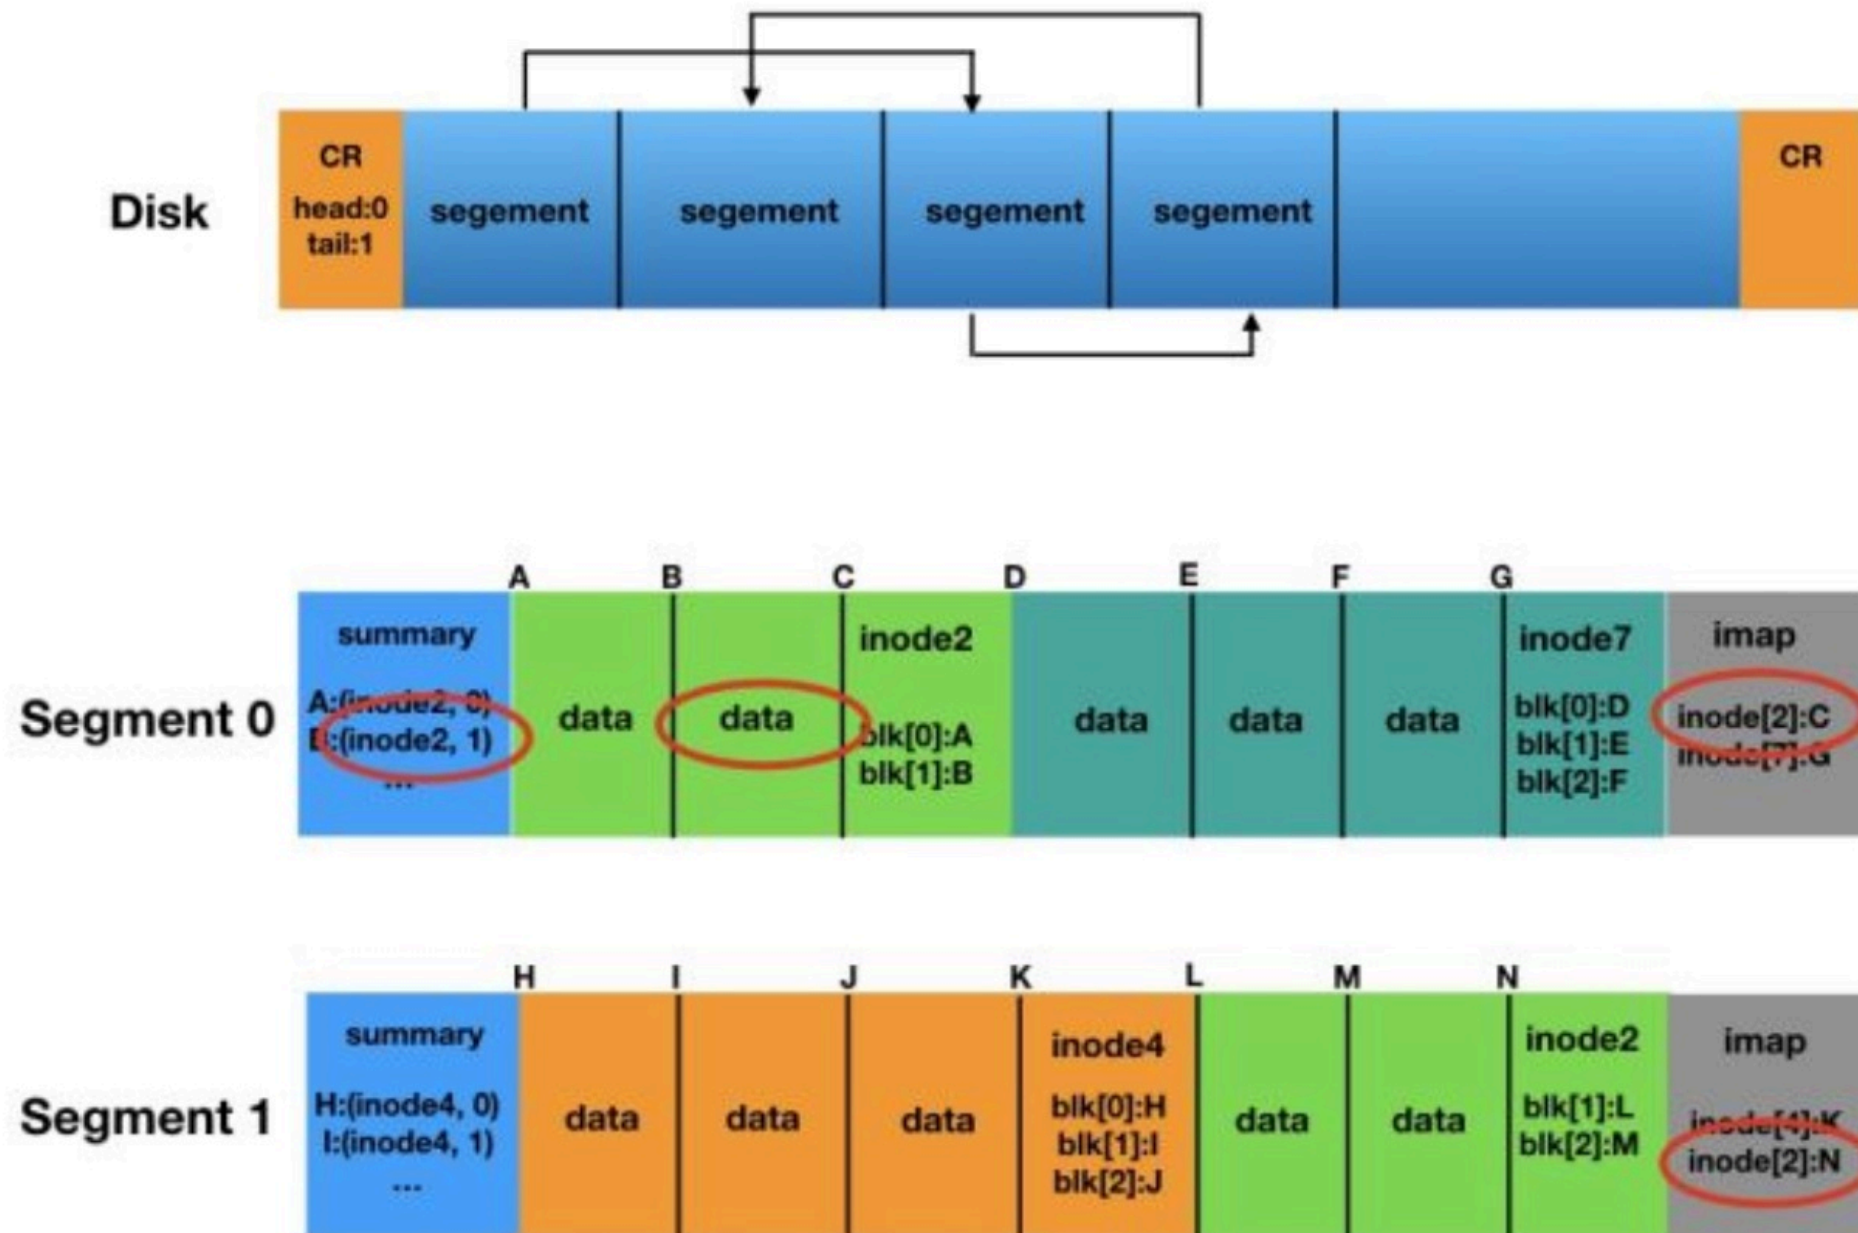
Directory is treated similar with file



Read

- First read CR
 - CR contains all the pointers to imap
- Read and cache imap
- Then given an inode number of a file, it looks up the imap to get the address of the data on the block
- Read data from block

Crash Recovery





Crash Recovery

- CR contains pointers to the head and tail segments
- Each segment points to next segment
- CR is updated periodically, 30s for example
- Segment is written into disk when it is full
- Crash can happen
 - Write to a segment
 - Write to the CR



Write to CR

- LFS keeps to CRs, and write to them alternatively
- Write protocol
 - First writes out header (with timestamp), then body, and last the one last block (with timestamp)
- If crash happen when writing CR, LFS can detect this by detecting the inconsistent of the timestamps,
- LFS always chooses to use the most recent CRT with consistent timestamps



Write to a segment

- If crash happens, then the CR has not been written into disk
- roll forward
 - Start with the last checkpoint, and find the end of the log, and then use that to find next segment and see if there are any new updates
 - Use this to recovery the data